

Dynamical systems time series classification

Kirill Semkin^{1*} and Vadim Strijov^{2*}

^{1*}Moscow Russia.

^{2*}Moscow Russia.

*Corresponding author(s). E-mail(s): semkin.ki32@gmail.com;
strijov@forecsys.ru;

Keywords: time series; classification; dynamical systems; NeuralODE; inertial measurement unit

We propose a method for multivariate time series classification under the assumption that series of different classes are generated by distinct dynamical systems, modeled as ordinary or stochastic differential equations whose noisy phase trajectories form the observed data. The classifier assigns each series to the class whose dynamics best fit it, and we prove it is asymptotically optimal under suitable noise conditions. Since exact trajectories are rarely available, we replace them with conditional expectations obtained by solving a smoothing problem, and learn unknown vector fields from data via Neural Ordinary Differential Equations. The method is validated on synthetic linear systems and human activity recognition from inertial sensors.

1 Introduction

The paper introduces a method for time series classification. A *time series* is a finite sequence of scalar values $x(t_i) = \{x_{t_1}, \dots, x_{t_n}\}$ of variable length defined at particular

times $t_1, \dots, t_n \in \mathbb{R}$. A *multivariate* time series is a sequence of vector values $\mathbf{x}(t_i) = \{\mathbf{x}_{t_1}, \dots, \mathbf{x}_{t_n}\}$, where $\mathbf{x}_{t_i} \in \mathbb{R}^d$. This is equivalent to having d time series with the same times. Everywhere further we assume time series to be multivariate in general. From a theoretical point of view it is convenient to assume the time series values $\mathbf{x}_{t_i}$ to be *random vectors*. Time series have an associated *class label* $y \in 1 \dots K$ that is also random in general. Here K is the number of classes. A *classifier* is a map $\delta(\mathbf{x}(t_i))$ that associates any time series with a class label. We will seek to obtain a “good” classifier in some sense defined later.

There are many methods for time series classification in the literature. *Distance-based* approaches rely on measuring similarity between time series using distance metrics. The metric can be defined using Dynamic Time Warping [1], and then the k-Nearest-Neighbors (k-NN) classifier can be used [2]. *Feature-based* approaches extract a finite set of statistical [3, 4] or spectral [5] features from time series for further classification. *Deep learning* approaches use neural networks to automatically learn time series features from raw training data. Convolutional neural networks [6, 7], recurrent neural networks (RNN) [8] or their combination [9–11] are designed to process sequential data like time series.

In the paper we assume a particular model behind the time series — a *dynamical system*. By that we mean the time-invariant *ordinary differential equation (ODE)* of the form

$$\begin{cases} \frac{d}{dt} \bar{\mathbf{x}}(t) = f(\bar{\mathbf{x}}), \\ \bar{\mathbf{x}}(0) = \bar{\mathbf{x}}_0, \end{cases} \quad (1)$$

where $\bar{\mathbf{x}} \in M$, $f : M \rightarrow M$ is a *vector field*, $\bar{\mathbf{x}}_0$ is a *starting point*, and M is a differentiable manifold in general. We assume the vector fields are not arbitrary but belong to some normed space $f \in (\mathcal{F}(M), \|\cdot\|)$. To find a solution $\bar{\mathbf{x}}(t)$ satisfying (1) on some time interval $t \in [0, T]$ is to solve an *initial value problem (IVP)*. We also denote the IVP solution as $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f)$ and call it a *phase trajectory*. Under some regularity

conditions on f the solution is unique [12]. We associate a different vector field f_j , $j \in 1, \dots, K$, with each series class. The series values are the noisy phase trajectories

$$\mathbf{x}(t_i) = \bar{\mathbf{x}}(t_i) + \boldsymbol{\epsilon}(t_i), \quad (2)$$

where $\boldsymbol{\epsilon}(t_i)$ is a random noise process. Its values $\boldsymbol{\epsilon}_{t_i}$ are i.i.d. random vectors (mutually for all classes and any times) having

$$\mathbb{E}[\boldsymbol{\epsilon}_{t_i}] = 0, \text{ Cov}(\boldsymbol{\epsilon}_{t_i}) = R \text{ where } \|R\|_2 \leq \sigma^2 < \infty. \quad (3)$$

Here R is the covariance matrix and $\|R\|_2$ is the matrix spectral norm. Therefore $\mathbb{E}[\mathbf{x}(t_i)] = \bar{\mathbf{x}}(t_i)$ and $\text{Cov}[\mathbf{x}(t_i)] = R$. In general, the classifier must account for both the deterministic component $\bar{\mathbf{x}}(t_i)$ and the noise component $\boldsymbol{\epsilon}(t_i)$ of the series. We will later show that under some conditions there is an asymptotically accurate classifier that does not rely on the noise. Then we will discard these conditions and introduce a more general classifier. Moreover, we will be able to adapt the classifier to the more general dynamical system described by the *stochastic differential equation (SDE)* [13] of the form

$$\begin{cases} d\bar{\mathbf{x}}(t) = f(\bar{\mathbf{x}})dt + GdW(t), \\ \bar{\mathbf{x}}(0) = \bar{\mathbf{x}}_0, \end{cases} \quad (4)$$

where $W(t)$ is an m -dimensional Brownian motion and G is a $d \times m$ constant diffusion matrix, unique for each class. Phase trajectories $\bar{\mathbf{x}}(t)$ are *random processes* in this model. The definition of the IVP translates straightforwardly for the SDE. The IVP solution is unique under some regularity conditions on f [13].

Our approach can be extended to a less direct connection between observed time series and the phase trajectories (2). Imagine the connection is given by some *observation*

map $\alpha : M \rightarrow \mathbb{R}^d$ such that

$$\mathbf{x}(t_i) = \alpha(\bar{\mathbf{x}}(t_i)) + \epsilon(t_i), \quad (5)$$

where M has dimension $k \neq d$ in general. We would like to reconstruct the $\bar{\mathbf{x}}(t_i)$ from the $\mathbf{x}(t_i)$ to build a classifier. The reconstruction will be based on building an *embedding* of the manifold M . In general, an embedding is a smooth map $\phi : M \rightarrow N$ into another manifold N that is a diffeomorphism onto its image $\phi(M)$. It follows that the phase trajectories $\bar{\mathbf{x}}(t) \subset M$ and $\phi(\bar{\mathbf{x}}(t)) \subset \phi(M)$ are basically identical (more strictly, diffeomorphic). To reconstruct the phase trajectories we need to build $\phi(\bar{\mathbf{x}}(t_i))$ from the observed time series $\mathbf{x}(t_i)$. One way to achieve this is the Whitney Embedding Theorem [14]. If α is a generic map from M to $\mathbb{R}^{2k+1}$ then α is itself an embedding, $\alpha = \phi$, and therefore $\mathbf{x}(t_i) = \phi(\bar{\mathbf{x}}(t_i))$. This is basically the original connection (2) but with the diffeomorphic phase trajectories instead of the original ones. The map is generic if it preserves information about initial dynamics. Usually it means that α is smooth and non-constant. However, this approach requires an excessive number, $2k + 1$, of time series even when k is rather small. Another approach is the Takens embedding theorem [15]. It requires only one observed scalar time series $x(t_i)$. If α is a generic map from M to $\mathbb{R}$ then the embedding is given by the map

$$\phi(\bar{\mathbf{x}}) = (\alpha(\bar{\mathbf{x}}), \alpha(f(\bar{\mathbf{x}})), \dots, \alpha(f^m(\bar{\mathbf{x}}))),$$

where $m > 2k$. Substituting $\bar{\mathbf{x}}(t_i)$ gives us $\phi(\bar{\mathbf{x}}(t_i)) = (x(t_i), x(t_{i-1}), \dots, x(t_{i-m}))$, which we call a *delay vector*. An illustration of the Takens theorem is in Fig. 1. The delay vectors form a noisy diffeomorphic phase trajectory because $x(t_i)$ is itself noisy. So they are basically identical to $\mathbf{x}(t_i)$ from (2). Further, we do not differentiate between reconstructed phase trajectories and the original ones.




Figure 1 Application of the Takens theorem. On the left are observations $x(t_i)$. To build a point of the phase trajectory $\bar{\mathbf{x}}(t_i)$ we must take several lagged observation values (colored in red), e.g. $\bar{\mathbf{x}}(t_i) = (x(t_i), x(t_{i-1}), x(t_{i-2}))^T$.

Our time series model can be applied to many real-world classification problems. For example, in medicine an electroencephalogram (EEG) is a direct, macroscopic measure of the brain’s electrical activity [16]. The activity is governed by Maxwell’s equations, which can be reduced to ODEs [17]. The electrocardiogram (ECG) can be used for detecting arrhythmias or other heart conditions [6], while the EEG can identify seizures and sleep stages from brainwave data [7]. In weather forecasting the fundamental equations of atmospheric dynamics can also be reduced to an ODE system [18]. The classification problems in this field include weather regime classification [19] and climate pattern identification [20]. In human activity recognition (HAR) ODEs are primarily used to model the dynamics of human motion. The core idea is that human movement is a continuous physical process governed by the laws of mechanics, which can be described by differential equations [21]. This idea can be applied to basic or fine-grained activity recognition [22, 23], fall detection [24], and other tasks. To conclude, our model is applicable when it is reasonable to assume that the observed time series come from some dynamical system of the form (1) or (4).

Our classifier will rely on the knowledge of the vector fields f_j , which is not always available. To compensate for this we will build approximate fields from data

using Neural Ordinary Differential Equations (NODE) [25]. It allows using gradient optimization when the optimized objective depends on ODE trajectories. There are plenty of papers using NODE for regression tasks. For example, NODE is used in the context of physics-informed ODEs [26, 27] and control learning [28, 29]. However, very few papers analyze NODE in the context of time series classification. The authors of [30] reinterpreted convolutional neural networks in terms of latent ODE dynamics and used them for hyperspectral image classification. The paper [31] solved the classification problem in terms of control theory and NODE. Our paper contributes to NODE applications for general dynamical systems.

The NODE paper [25] introduced its own popular time series model called latent ODE. It is illustrated in Fig. 2. The model is probabilistic and in our terms it can be described as

$$\begin{aligned}\bar{\mathbf{x}}_{t_0} &\sim p_0(\bar{\mathbf{x}}_{t_0}), \\ \bar{\mathbf{x}}_{t_1}, \bar{\mathbf{x}}_{t_2}, \dots, \bar{\mathbf{x}}_{t_N} &= \text{ODESolve}(\bar{\mathbf{x}}_{t_0}, f, t_0, \dots, t_N), \\ \text{each } \mathbf{x}_{t_i} &\sim p_x(\mathbf{x}_{t_i} | \bar{\mathbf{x}}_{t_i})\end{aligned}$$

where $\text{ODESolve}()$ is an approximate solution of the initial value problem using an ODE solver [32]. The authors used a variational autoencoder (VAE) [33] to infer the distributions p_0, p_x and the vector field f from the $\mathbf{x}_{t_i}$. While this model is more general and is widely used [34–36], it has several drawbacks compared to ours. First, it has to parameterize f, p_x and the VAE. Hence, more data is needed in general to obtain good approximations of these functions. Our approach may only need to parameterize f . Second, in the latent ODE model the connection between the phase trajectories and the observed time series is given by the general distribution p_x . Thus, the connection is difficult to interpret. In addition, the VAE is needed to reconstruct

the $\bar{\mathbf{x}}(t_i)$ from the $\mathbf{x}(t_i)$. Our reconstruction approaches discussed above give explicit and easy ways to obtain $\bar{\mathbf{x}}(t_i)$ from the $\mathbf{x}(t_i)$ and vice versa. Several papers can be found in the literature [37, 38] that use Takens theorem ideas with NODE to restore latent dynamics. Unfortunately, this is poorly covered for machine learning applications. Our paper aims to close this gap.




Figure 2 Computational graph of the latent ODE model. Source: reprinted from Chen et al. (2018). Here $z(t_i)$ corresponds to $\bar{\mathbf{x}}(t_i)$.

In the next section we will describe a general principle for the time series classification inside our model. We will build an asymptotically optimal classifier under the mentioned properties of the noise ϵ_{t_i} and several additional conditions. These conditions will help with the theoretical proof but may not hold in practice. We will then progressively discard these conditions and adapt the classifier to more general settings. We will also show how to adapt the classifier to the SDE case (4). Finally, we will describe the procedure of building vector field approximations from the data. Then we will apply our method to two classification tasks for empirical validation. The first will be the classification of three linear ODE systems. The second will be the classification of human motions from Inertial Measurement Unit (IMU) sensors.

2 Asymptotically optimal classifier

The main principle for building a classifier in the paper is the proximity of an observed time series $\mathbf{x}(t_i)$ to the theoretical initial value problem solution $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f_j)$ for each class $j \in 1 \dots K$. Basically, our classifier is distance-based.

Denote by $\rho_n(\mathbf{x}(t_i), \mathbf{z}(t_i))$ the metric between time series $\mathbf{x}(t_i)$, $\mathbf{z}(t_i)$ of length n defined as

$$\rho_n(\mathbf{x}(t_i), \mathbf{z}(t_i)) = \sum_{i=1}^n \|\mathbf{x}(t_i) - \mathbf{z}(t_i)\|_2^2. \quad (6)$$

Then, define the classifier as

$$\delta(\mathbf{x}(t_i)) = \arg \min_{j \in 1 \dots K} \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)). \quad (7)$$

It turns out that this theoretical classifier has a vanishing probability of misclassification as the length of the time series increases. But the starting point $\bar{\mathbf{x}}_0$ must be known exactly and the phase trajectories from different systems must diverge in some sense. The formal result is as follows:

Theorem 1. *Let $\mathbf{x}(t_i) = \{\mathbf{x}_{t_1}, \dots, \mathbf{x}_{t_n}\}$ be the time series to classify, defined for all $n \in \mathbb{N}$. Let $\bar{\mathbf{x}}_0$ be the starting point of the associated phase trajectory $\bar{\mathbf{x}}(t_i)$. Moreover, let*

$$\rho_\infty(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) = +\infty \quad \forall j \neq p; \quad j, p \in 1 \dots K.$$

Denote by y the true class of the time series. Then, under proposed conditions on the noise (3), the following holds

$$\mathbb{P}(\delta(\mathbf{x}(t_i)) \neq y) \rightarrow 0, \quad n \rightarrow \infty.$$

The proof of the theorem is deferred to [Appendix A](#).

3 Why optimal classifier may be intractable

Unfortunately, the classifier (7) relies on precise IVP solutions and Theorem 1 relies on a precise starting point $\bar{\mathbf{x}}_0$. These requirements are often too strong in practice. We now provide several reasons for that and propose a new classifier.

3.1 Approximate vector fields

Imagine instead of the true vector fields f_j we only have approximate versions $\tilde{f}_j$ such that

$$\|f_j - \tilde{f}_j\| \leq r_j,$$

where $f_j, \tilde{f}_j \in \mathcal{F}(M)$ and r_j are approximation error bounds. This situation always arises when one approximates vector fields from the data (see Section 6).

In general, an arbitrarily small approximation error is enough for the phase trajectories of f and $\tilde{f}$ to diverge. We show such a system in Example 1 in Appendix A. The visualization of the associated trajectories is in Fig. 3. That makes classifier (7) useless. Nonetheless, for a large family of vector fields the $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, \tilde{f}_j)$ are arbitrarily close to $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f_j)$ if r_j is sufficiently small and the times are bounded, $t < T$. These are called Grönwall-type bounds and can be found in the ODE literature [39]. So further we assume the approximation errors are small enough that we do not differentiate between $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, \tilde{f}_j)$ and $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f_j)$. We also assume that the vector fields f_j are considerably different across classes to exclude situations like in Example 1.

3.2 Approximate IVP solutions and noisy phase trajectories

Even if the vector fields are known exactly, the IVP is almost always intractable [32]. The numerical methods called ODE solvers are able to approximate the IVP solution with arbitrary precision [32] for regular vector fields. The main approximation




Figure 3 Diverging phase trajectories of the close vector fields $\|f - \tilde{f}\| \leq 0.1$ with $\bar{\mathbf{x}}_0 = (0, 1)^T$. While $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f)$ is a circle trajectory (called a center in dynamical systems), the $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, \tilde{f})$ is a spiral sink that converges to $\mathbf{0}$.

parameter of the solvers is the integration step τ . A solver has order of convergence l if

$$\sup_{t \in [0, T]} |\text{ODESolve}(t|\bar{\mathbf{x}}_0, f) - \bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f)| = \mathcal{O}(T\tau^l). \quad (8)$$

The higher the convergence order and the smaller τ , the higher the computational complexity of the ODESolve procedure. Besides this tradeoff, some vector fields imply additional constraints on the smallness of τ . If these are not satisfied, the solver solution diverges from the true one. This effect is called *stiffness* and is usually addressed by choosing an appropriate solver [40]. The illustration of this effect is in Fig. 4.

As follows from (8), applying ODESolve to find the full trajectory $t \in [0, T]$ starting from the exact $\bar{\mathbf{x}}_0$ becomes impractical for long trajectories. We also assume

$$\bar{\mathbf{x}}_0 \sim F(\cdot) \quad (9)$$




Figure 4 Exact and approximate IVP solutions for the stiff system $f(\bar{\mathbf{x}}) = \begin{pmatrix} -35 & 0 \\ 0 & 0.5 \end{pmatrix} \bar{\mathbf{x}}$ and $\bar{\mathbf{x}}_0 = (1000, 1)^T$. The Runge-Kutta23 (RK23) ODE solver approximates the exact solution with high precision. The Euler (euler) ODE solver is unstable and diverges from the exact solution. The integration step $\tau = 0.05$ is the same for both solvers, but for Euler it is not small enough. The RK23 has a higher convergence order and is stable under such τ .

to be a random vector in general, where $F(\cdot)$ is its distribution function. Hence, the classifier (7) cannot be applied under these conditions.

On the other hand, imagine we knew $\bar{\mathbf{x}}_0$ and $\bar{\mathbf{x}}_{t_i}$ for some $i \in 1 \dots n$. We could then approximate $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f), t \in [0, t_i)$ using $\text{ODESolve}(t|\bar{\mathbf{x}}_0, f)$, and approximate $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f), t \in [t_i, T]$ using $\text{ODESolve}(t|\bar{\mathbf{x}}_{t_i}, f)$. The solver's accuracy increases. To obtain a practical solution, we could simply replace $\bar{\mathbf{x}}_{t_i}$ with $\mathbf{x}_{t_i}$. But if the σ from (3) is large enough, then the starting points differ considerably with high probability. This leads to very different IVP solutions $\text{ODESolve}(t|\mathbf{x}_{t_i}, f), \text{ODESolve}(t|\bar{\mathbf{x}}_{t_i}, f)$ in general [39]. We must take the randomness into account in a different way.

4 Classification as smoothing problem

Let us break the dynamics (1) into consecutive time segments $[t_0, t_1), [t_1, t_2), \dots, [t_{n-1}, t_n]$, so we have n IVPs

$$\begin{cases} \frac{d}{dt}\bar{\mathbf{x}}(t) = f(\bar{\mathbf{x}}), \\ \bar{\mathbf{x}}(0) = \bar{\mathbf{x}}_{t_i}, \\ t \in [0, t_{i+1} - t_i]. \end{cases} \quad (10)$$

Let us approximate each IVP solution by ODESolve, so that $\bar{\mathbf{x}}_{t_{i+1}} = \text{ODESolve}(t_{i+1} - t_i | \bar{\mathbf{x}}_{t_i}, f)$. Combining this with (2) and (9), we obtain the discrete dynamics

$$\begin{cases} \bar{\mathbf{x}}_{t_{i+1}} = \text{ODESolve}(t_{i+1} - t_i | \bar{\mathbf{x}}_{t_i}, f), \\ \mathbf{x}_{t_i} = \bar{\mathbf{x}}_{t_i} + \epsilon_{t_i}, \\ \bar{\mathbf{x}}_0 \sim F(\cdot). \end{cases} \quad (11)$$

The $\mathbf{x}(t_i)$ is observed, everything else is not. To adapt classifier (7) to this model we need to replace the intractable $\bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f)$ with the best approximation based on the random $\bar{\mathbf{x}}_0$ and $\mathbf{x}_{t_i}$, $i \in 1 \dots n$. If we take an L_2 norm in the space of random vectors, the best approximation is the *conditional expectation* [41]. We denote the series of conditional expectations by $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$, which means the i -th series value is $\mathbb{E}[\bar{\mathbf{x}}_{t_i} | \bar{\mathbf{x}}_0, \mathbf{x}_{t_1}, \dots, \mathbf{x}_{t_n}]$. The problem of finding $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$ is called the *smoothing problem*. On the one hand, it is intractable in general, similarly to the IVP. On the other hand, there are numerical methods for obtaining approximate solutions. Some of them are extensions of the Kalman filter [42, 43]. Others are based on Bayesian and particle filters [44].

The new classifier is defined as

$$\delta(\mathbf{x}(t_i)) = \arg \min_{j \in 1 \dots K} \rho_n(\mathbf{x}(t_i), \mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f_j]). \quad (12)$$

The classifier now relies not only on the dynamics driven by the vector fields f_j but also on the observations $\mathbf{x}(t_i)$. This mitigates the problems discussed in Section 3:

1. Inaccuracy from the approximate $\tilde{f}$ and ODEsolve decreases because we give up the full IVP $t \in [0, T]$ in favor of the smaller IVPs on time segments $t \in [t_i, t_{i+1}]$.
2. Randomness of the $\bar{\mathbf{x}}_0, \mathbf{x}_{t_i}$ is taken into account via the conditional expectation $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$. Informally, it optimally balances the underlying dynamics (1) with the noisy observations to approximate the IVP solution $\bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f)$.

Comparison of the IVP $\bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f)$ approximation by ODEsolve($t_i | \bar{\mathbf{x}}_0, f$) and by $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$ is illustrated in Fig. 5.

5 Adapting classifier for SDE

We can segment the dynamics (4) similarly to (10) and obtain

$$\begin{cases} d\bar{\mathbf{x}}(t) = f(\bar{\mathbf{x}})dt + GdW(t), \\ \bar{\mathbf{x}}(0) = \bar{\mathbf{x}}_{t_i}, \\ t \in [0, t_{i+1} - t_i]. \end{cases} \quad (13)$$

If all time segments are small enough, we can approximate the dynamics $f(\bar{\mathbf{x}})dt \approx \text{ODEsolve}(t_{i+1} - t_i | \bar{\mathbf{x}}_{t_i}, f)$ and $GdW(t) \approx G \cdot \mathcal{N}(0, (t_{i+1} - t_i) \cdot I) = \mathcal{N}(0, (t_{i+1} - t_i) \cdot$




Figure 5 Different IVP approximations for the one-dimensional system $f(x) = -18x$. The initial value $x_0 \sim \mathcal{N}(1, 1)$ and $\epsilon_{t_i} \sim \mathcal{N}(0, 1)$. ODESolve uses Euler solver with initial value $\mathbb{E}[x_0] = 1$. Stepwise ODESolve uses the Euler solver on consecutive time segments $[t_i, t_{i+1})$ with noisy initial values x_{t_i} . The Unscented Kalman Filter (UKF) [43] is used to compute $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$. UKF uses stepwise Euler internally. We can see that UKF demonstrates the best IVP approximation by using both the underlying dynamics and the noisy observations. UKF is the most accurate IVP approximation with probability 0.92 by Monte-Carlo estimation.

$GG^T) := \xi_{t_{i+1}}$. As a result, we obtain discrete dynamics akin to (11):

$$\begin{cases} \bar{\mathbf{x}}_{t_{i+1}} = \text{ODESolve}(t_{i+1} - t_i | \bar{\mathbf{x}}_{t_i}, f) + \xi_{t_{i+1}}, \\ \mathbf{x}_{t_i} = \bar{\mathbf{x}}_{t_i} + \epsilon_{t_i}, \\ \bar{\mathbf{x}}_0 \sim F(\cdot). \end{cases} \quad (14)$$

Finding $\mathbb{E}[\bar{\mathbf{x}}(t_i) | \bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$ here is just a more general smoothing problem and can be solved by the same methods listed above. The classifier (12) for the SDE case stays the same. Comparison of the different SDE IVP $\bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f)$ approximations is given in Fig. 6.




Figure 6 Different IVP approximations for the one-dimensional system $dx(t) = xdt + dW(t)$. The initial value $x_0 \sim \mathcal{N}(1, 1)$ and $\epsilon_{t_i} \sim \mathcal{N}(0, 1)$. The phase trajectory is a random process here, in contrast to the ODE dynamics. ODEsolve uses the Euler solver with initial value $\mathbb{E}[x_0] = 1$. We can see it is an inadequate model for the SDE dynamics. Stepwise ODEsolve uses the Euler solver on consecutive time segments $[t_i, t_{i+1})$ with noisy initial values x_{t_i} . It is more accurate but is subject to the noise and still misses the $dW(t)$ term. The Unscented Kalman Filter (UKF) is used to compute $\mathbb{E}[\bar{x}(t_i) | \bar{x}_0, \mathbf{x}(t_i), f]$. UKF uses stepwise Euler internally but also accounts for the $dW(t)$ term. It is the most accurate IVP approximation among those mentioned with probability 0.97 by Monte-Carlo estimation.

6 Learning vector field from data

Sometimes the vector field f may be unknown. Instead, time series samples from the associated dynamics (1) may be available. From (1) and (2) we have

$$\mathbf{x}_{t_i} = \bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f) + \epsilon_{t_i}, \quad i \in 1 \dots n,$$

where $\bar{\mathbf{x}}_0$ is defined as in (9). Approximating the true IVP by ODEsolve, we obtain

$$\mathbf{x}_{t_i} = \text{ODEsolve}(t_i | \bar{\mathbf{x}}_0, f) + \epsilon_{t_i}, \quad i \in 1 \dots n.$$

By taking the expectation of both sides and gathering the terms on the left-hand side, we get

$$\mathbb{E}[\mathbf{x}_{t_i} - \text{ODESolve}(t_i|\bar{\mathbf{x}}_0, f)] = 0, \quad i \in 1 \dots n.$$

We can apply the *generalized method of moments* (GMM) [45] here. Assume we have a set of vector fields $\mathcal{S}_\Theta \subset \mathcal{F}(M)$ parameterized by the real vector Θ . Assume $f \in \mathcal{S}_\Theta$. Denote the solution of the following optimization problem

$$\sum_{i=1}^n \|\mathbf{x}_{t_i} - \text{ODESolve}(t_i|\bar{\mathbf{x}}_0, f)\|_2^2 \rightarrow \min_{f \in \mathcal{S}_\Theta} \quad (15)$$

as $f_{\hat{\Theta}}$. Here $\bar{\mathbf{x}}_0, \mathbf{x}_{t_i}, i \in 1 \dots n$, are sampled values. If some technical conditions on Θ hold [45], the GMM guarantees that

$$f_{\hat{\Theta}} \xrightarrow{\mathbb{P}} f, \quad n \rightarrow \infty. \quad (16)$$

Problem (15) can be solved using NODE. As discussed in Section 3, n should not be large due to approximation errors of the ODESolve. Instead, it is better to sample N short batches $B_j = (\bar{\mathbf{x}}_0, \{\mathbf{x}_{t_i}\}_{i=1}^n)$ and solve

$$\sum_{j=1}^N \sum_{\bar{\mathbf{x}}_0, \mathbf{x}_{t_i} \in B_j} \|\mathbf{x}_{t_i} - \text{ODESolve}(t_i|\bar{\mathbf{x}}_0, f)\|_2^2 \rightarrow \min_{f \in \mathcal{S}_\Theta}, \quad (17)$$

which is equivalent to (15) in terms of the convergence (16).

SDE dynamics (4) can also be learned from data by methods similar to NODE [46, 47].

7 Empirical study

In this section we validate the proposed classifier (12) on two classification tasks with progressively decreasing prior knowledge about the underlying dynamics. In the first

task the vector fields f_j are known analytically, which lets us isolate the behaviour of the smoothing-based classifier from the error of the field approximation. In the second task the fields are unknown and are learned from data with NODE as described in Section 6. The phase trajectories are still observed directly through a multivariate IMU sensor. Data is taken from the [MotionSense](#) dataset [48].

7.1 Experimental setup

Across all tasks we evaluate the classifier (12) where the conditional expectation $\mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f_j]$ is approximated with the Unscented Kalman Filter together with the Rauch–Tung–Striebel smoother [44]. The smoother uses an Euler solver for the ODEsolve in the (11).

Classification quality is reported as accuracy on a held-out test set. All trajectory integration and vector field learning are performed with `torchdiffeq`, and the trajectory smoothing with `filterpy`.

7.2 Linear dynamical systems

The first task classifies trajectories generated by three linear ODE systems with the following form

$$f(\bar{\mathbf{x}}) = A\bar{\mathbf{x}} = U\Lambda U^T\bar{\mathbf{x}} = U \begin{pmatrix} -\gamma_1 & w & 0 \\ -w & -\gamma_1 & 0 \\ 0 & 0 & -\gamma_2 \end{pmatrix} U^T\bar{\mathbf{x}}. \quad (18)$$

Here A is a real matrix with the specified block decomposition, where U is orthogonal. All three systems share the same Λ but differ in the matrix U . The vector field (18) simulates damped oscillation along the first two basis directions and exponential decay along the third. The basis directions are the U columns. We used $\gamma_1 = 0.7, \gamma_2 = 0.6, w = 4$. For each class we sample $\bar{\mathbf{x}}_0 \sim \mathcal{N}(\mathbf{0}, 4 \cdot I)$ and integrate the

system trajectories of length $n = 100$ with step $\tau = 0.05$. Then we corrupt them with additive Gaussian noise of level $\sigma = 2$ following (2)–(3). The visualization of the final series and the underlying phase trajectories is in Fig. 7 (with a lower noise level for clearer illustration).




Figure 7 Observed time series $\mathbf{x}(t_i)$ generated from each linear ODE system (y label in the figure legend) and smoothed phase trajectories from every field f_0, f_1, f_2 . The systems are a combination of exponential decay and damped oscillations. Noise level for the illustration is $\sigma = 0.2$.

Because the true fields f_j are available, no field learning is required. We compare our classifier against time series classification baselines: a k -Nearest-Neighbours classifier with the Dynamic Time Warping distance [2] and a recurrent (LSTM) network trained end-to-end on the raw series. Both baselines are relevant to our model. Recurrent networks have discrete dynamics inside that can be seen as Euler solver steps [25]. DTW is used to define a distance between time series [2] different from (6). Training

curves for the LSTM are presented in Fig. 8. Classification accuracy is reported in Table 1.




Figure 8 Learning curves for the LSTM training: MSE loss on train and validation series.

Table 1 Classification accuracy on the linear ODE systems task ($K = 3$, $n = 100$, $\sigma = 2$).

Method	Accuracy
k -NN + DTW	0.67
LSTM	0.27
Ours (UKF smoothing)	0.86

From Fig. 8 we can see that the LSTM was not able to learn the ODE systems. The model cannot differentiate them due to the high noise level and the systems’ “proximity” implied by the shared Λ . The k -NN + DTW classifier is distance-based like (12) but is sub-optimal. It uses DTW to compute the series distance instead of (6) and $\mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]$. It also learns f_j only indirectly by memorizing the training series for each class. Classifier (12) combines information from both f_j and the noise level and shows the best performance, though it is not absolutely precise due to the high noise level and the limited trajectory length.

7.3 Human motion recognition from IMU

The second task uses the BasicMotions dataset of Inertial Measurement Unit (IMU) recordings. The dataset contains the rotation rate (w_x, w_y, w_z) and acceleration (a_x, a_y, a_z) from the IMU (located on the wrist) sampled at 10 Hz. Subjects performed 4 activities: standing, walking, running, and badminton, repeating each 5 times. Each activity corresponds to a classification class. The train and test splits contain 2 subjects that makes 40 six-dimensional time series per split. The series' length is 100 measurements.

The human body in motion is well approximated as a system of articulated rigid bodies (limbs and trunk coupled at joints), and the motion of such a system is governed by the Lagrangian equations of mechanics [49, 50]. These equations can be written in the form (1). The IMU measures only part of the full system, but we further show that it is enough for classification. Steady locomotion activities (walking, running, badminton) are rhythmic and quasi-periodic, which in dynamical-systems language corresponds to stable limit cycles. Static postures (standing) correspond to stable equilibria.

We apply both ODE and SDE approaches for this task. Unlike the previous experiment, the vector fields/diffusion terms for the activities are unknown. In ODE case, for each class we learn vector fields approximations from short trajectory batches by solving the objective (17). In SDE case, we discretised the dynamics (4) by the Euler-Maruyama scheme [?] that implies the Gaussian transition

$$\bar{\mathbf{x}}_{t_{i+1}} | \bar{\mathbf{x}}_{t_i} \sim N(\bar{\mathbf{x}}_{t_{i+1}} + f(\bar{\mathbf{x}}_{t_{i+1}})(t_{i+1} - t_i), (t_{i+1} - t_i) \cdot GG^T).$$

The fields f_j and the diffusion matrices G_j are then learned by maximizing likelihood of the train trajectories. In both approaches, the vector fields were parameterized by the neural nets.

We assume (2) as a model for the IMU time series. We take the rotation rate as phase trajectory $\mathbf{x}(t)$ because its derivative is directly connected to rigid body mechanics. Acceleration is itself a derivative of the velocity and shouldn't be included in the ODE/SDE dynamics. However, it indirectly contains properties of the motion, for example, speed when running or walking. So we integrate acceleration in the dynamics as a *control* variable — we let the vector field and diffusion matrices to depend on it. For example, ODE dynamics now writes as

$$\frac{d}{dt}(\bar{w}_x(t), \bar{w}_y(t), \bar{w}_z(t))^T = f(\bar{w}_x(t), \bar{w}_y(t), \bar{w}_z(t), \bar{a}_x(t), \bar{a}_y(t), \bar{a}_z(t))$$

We also modified classifier (12) by adding $\text{Cov}(\mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f])^{-1}$ weighting in the $\rho_n(\cdot, \cdot)$:

$$\begin{aligned} \rho_n(\mathbf{x}(t_i), \mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]) &= \sum_{i=1}^n (\mathbf{x}(t_i) - \mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f])^T \cdot \\ &\cdot \text{Cov}(\mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f])^{-1} (\mathbf{x}(t_i) - \mathbb{E}[\bar{\mathbf{x}}(t_i)|\bar{\mathbf{x}}_0, \mathbf{x}(t_i), f]). \end{aligned} \quad (19)$$

It allows to account for the uncertainty of the smoother and rely less on the high variance predictions. It was crucial heuristic for two reasons. First, the trajectories were rather short and we couldn't rely on asymptotic properties of the original classifier (Section 2). Second, human motions have natural, unavoidable variations called motor variability as well as sudden deviations in dynamics, see Fig. ???. These factors make smoothing highly variant and can't be accounted in the classifier without adding some measure of uncertainty.

The second task uses the BasicMotions dataset of Inertial Measurement Unit (IMU) recordings. We took the rotation rate (x, y, z) and acceleration (x, y, z) from the

IMU sampled at 10Hz as our time series. We assume them to be noisy phase trajectories following (2)–(3). We chose 6 subjects who performed 6 activities: walking, jogging, walking upstairs, walking downstairs, standing and sitting. Each activity corresponds to a classification class. Unlike the previous experiment, the vector fields for the activities are unknown. So for each class and each subject we learn an approximate field f_j from short trajectory batches by solving the objective (17). Subjects performed each activity at least twice, so we took one activity trial for learning and one trial for measuring accuracy. The noise level σ was estimated directly from the training data for each subject and activity.

Visualization of the time series and smoothed trajectories corresponding to different activities is in Fig. 9. We can separate the smoothed trajectories for the activities, which means the classes have different vector fields. Accuracy is reported in Table 2. High values imply the ODE model truly suits human motion data. The binomial test’s p-values for the individuality hypothesis are reported in Table 3. We also included the subject classification accuracy here using (??). The null hypothesis was rejected for jogging and walking and was not rejected for the other activities. This seems reasonable because sitting and standing are static postures with no articulated movements, in contrast to jogging and walking. The absence of individuality for walking up/downstairs was not expected and suggests an interesting insight.

Table 2 Activity classification accuracy on MotionSense IMU data.

Activity	Accuracy
jog	0.84
walk	0.84
upstairs	0.72
downstairs	0.70
stand	0.75
sit	0.74
all	0.765




Figure 9 The upper row shows rotation rate and acceleration time series for the jogging activity. Smoothed trajectories from f_{jog} and f_{wlk} are visualized for comparison. The lower row depicts time series for the upstairs activity with f_{ups} and f_{sit} smoothed trajectories. All series and fields were taken from subject 1.

Table 3 Subject classification p-values and accuracy across activities.

Activity	p-value	Accuracy
jog	0.03	0.84
walk	0.01	1.00
upstairs	0.35	0.68
downstairs	0.34	0.71
stand	0.41	0.65
sit	0.89	0.34
all	-	0.7

8 Conclusion

In the paper we have introduced a distance-based time series classifier. It assumes a particular time series model — ODE or SDE dynamical systems. We proved it

is asymptotically accurate under idealized conditions and reasonably adapted it to the more general case. Our classifier is simple and powerful as it requires only the systems' vector fields f_j . It then uses the well-developed framework of ODE solvers and smoothing methods. Our classifier is also flexible. If we do not know the vector fields, we can learn them from data. If we do not observe the entire dynamics, we can use the Whitney/Takens theorems to restore it. The classifier's flexibility and applicability were demonstrated empirically on human motion signals.

Our classifier has certain limitations by design. Even though the introduced dynamical systems are quite general, they cannot describe any time series. For example, we expect poor performance for non-continuous dynamics or general random processes. Next, learning vector fields may be cumbersome if we have no clue about their form or lack data. Finally, the numerical methods the classifier uses have their own conditions and bounds that should be understood before application.

There are several directions to extend the current research. First, the classifier may be improved if the exact distributions of the noise and the initial value are known (e.g. Student or Uniform). The theoretical classifier could be derived and empirically validated. Second, the classifier may be adapted to more general SDE dynamics. Currently, the diffusion term is $GdW(t)$ because most smoothing methods require i.i.d. diffusion samples. The general term $g(x(t), t)dW(t)$ produces dependent samples with different distributions. Lastly, the classifier may be adapted to the *change point detection* task. The change point would correspond to shifting from one dynamical system to the other. However, this requires separate investigation.

References

- [1] Berndt, D.J., Clifford, J.: Using dynamic time warping to find patterns in time series. In: Proceedings of the 3rd International Conference on Knowledge Discovery and Data Mining. AAAIWS'94, pp. 359–370. AAAI Press, Seattle, WA (1994)

- [2] Ding, H., Trajcevski, G., Scheuermann, P., Wang, X., Keogh, E.: Querying and mining of time series data: experimental comparison of representations and distance measures. *Proc. VLDB Endow.* **1**(2), 1542–1552 (2008) <https://doi.org/10.14778/1454159.1454226>
- [3] Nanopoulos, A., Alcock, R., Manolopoulos, Y.: Feature-based classification of time-series data, pp. 49–61. Nova Science Publishers, Inc., USA (2001)
- [4] Wang, X., Smith, K., Hyndman, R.: Characteristic-based clustering for time series data. *Data Min. Knowl. Discov.* **13**(3), 335–364 (2006) <https://doi.org/10.1007/s10618-005-0039-x>
- [5] Mörchen, F.: Time series feature extraction for data mining using dwt and dft (2003)
- [6] Rajpurkar, P., Hannun, A., Haghighpanahi, M., Bourn, C., Ng, A.: Cardiologist-level arrhythmia detection with convolutional neural networks (2017) <https://doi.org/10.48550/arXiv.1707.01836>
- [7] Lawhern, V., Solon, A., Waytowich, N., Gordon, S., Hung, C., Lance, B.: Eegnet: a compact convolutional neural network for eeg-based brain–computer interfaces. *Journal of Neural Engineering* **15** (2018) <https://doi.org/10.1088/1741-2552/aace8c>
- [8] Malhotra, P., Ramakrishnan, A., Anand, G., Vig, L., Agarwal, P., Shroff, G.: Lstm-based encoder-decoder for multi-sensor anomaly detection (2016) <https://doi.org/10.48550/arXiv.1607.00148>
- [9] Trigeorgis, G., Ringeval, F., Brueckner, R., Marchi, E., Nicolaou, M.A., Schuller, B., Zafeiriou, S.: Adieu features? end-to-end speech emotion recognition using a deep convolutional recurrent network. In: 2016 IEEE International Conference

- on Acoustics, Speech and Signal Processing (ICASSP), pp. 5200–5204 (2016).
<https://doi.org/10.1109/ICASSP.2016.7472669>
- [10] Hammerla, N., Halloran, S., Ploetz, T.: Deep, convolutional, and recurrent models for human activity recognition using wearables. (2016)
- [11] Shi, X., Chen, Z., Wang, H., Yeung, D.-Y., Wong, W.K., WOO, W.-c.: Convolutional lstm network: A machine learning approach for precipitation nowcasting (2015)
- [12] Bifurcation of Fixed Points of Vector Fields, pp. 356–497. Springer, New York, NY (2003). https://doi.org/10.1007/0-387-21749-5_21 . https://doi.org/10.1007/0-387-21749-5_21
- [13] Øksendal, B.: Stochastic Differential Equations vol. 77, p. 326 (2000). <https://doi.org/10.2307/3619809>
- [14] Whitney, H.: Differentiable manifolds. *Annals of Mathematics* **37**(3), 645–680 (1936). Accessed 2025-11-09
- [15] Takens, F.: Detecting strange attractors in turbulence. In: *Dynamical Systems and Turbulence*, Warwick 1980: Proceedings of a Symposium Held at the University of Warwick 1979/80, pp. 366–381 (2006). Springer
- [16] Buzsáki, G., Anastassiou, C., Koch, C.: The origin of extracellular fields and currents—eeg, ecog, lfp and spikes. *Nature reviews. Neuroscience* **13**, 407–20 (2012)
<https://doi.org/10.1038/nrn3241>
- [17] Daunizeau, J., Friston, K.J., Kiebel, S.J.: Variational bayesian identification and prediction of stochastic nonlinear dynamic causal models. *Physica D: Nonlinear Phenomena* **238**(21), 2089–2118 (2009) <https://doi.org/10.1016/j.physd.2009.08.002>

- [18] Kalnay, E.: Atmospheric Modeling, Data Assimilation and Predictability. Cambridge University Press, Cambridge (2002)
- [19] Michelangeli, P.-A., Vautard, R., Legras, B.: Weather regimes: Recurrence and quasi stationarity. *J. Atmos. Sci.* **52**, 1237–1256 (1995) [https://doi.org/10.1175/1520-0469\(1995\)052<1237:WRRAS>2.0.CO;2](https://doi.org/10.1175/1520-0469(1995)052<1237:WRRAS>2.0.CO;2)
- [20] Wheeler, M., Hendon, H.: An all-season real-time multivariate mjo index: Development of an index for monitoring and prediction. *Monthly Weather Review - MON WEATHER REV* **132**, 1917–1932 (2004) [https://doi.org/10.1175/1520-0493\(2004\)132<1917:AARMMI>2.0.CO;2](https://doi.org/10.1175/1520-0493(2004)132<1917:AARMMI>2.0.CO;2)
- [21] Zhao, J., Li, X., Shum, C., McPhee, J.: A review of physics-based and data-driven models for real-time control of polymer electrolyte membrane fuel cells. *Energy and AI* **6**, 100114 (2021) <https://doi.org/10.1016/j.egyai.2021.100114>
- [22] Anguita, D., Ghio, A., Oneto, L., Parra, X., Reyes-Ortiz, J.L.: A public domain dataset for human activity recognition using smartphones. In: *The European Symposium on Artificial Neural Networks* (2013). <https://api.semanticscholar.org/CorpusID:6975432>
- [23] Mekruksavanich, S., Jitpattanakul, A.: Deep convolutional neural network with rnns for complex activity recognition using wrist-worn wearable sensor data. *Electronics* **10**, 1685 (2021) <https://doi.org/10.3390/electronics10141685>
- [24] Igual, R., Medrano, C., Plaza, I.: Challenges, issues and trends in fall detection systems. *Biomedical engineering online* **12**, 66 (2013) <https://doi.org/10.1186/1475-925X-12-66>
- [25] Chen, R.T.Q., Rubanova, Y., Bettencourt, J., Duvenaud, D.: Neural ordinary differential equations. In: *Proceedings of the 32nd International Conference*

- on Neural Information Processing Systems. NIPS'18, pp. 6572–6583. Curran Associates Inc., Red Hook, NY, USA (2018)
- [26] Lai, Z., Mylonas, C., Nagarajaiah, S., Chatzi, E.: Structural identification with physics-informed neural ordinary differential equations. *Journal of Sound and Vibration* **508**, 116196 (2021) <https://doi.org/10.1016/j.jsv.2021.116196>
 - [27] Lai, Z., Liu, W., Jian, X., Bacsá, K., Sun, L., Chatzi, E.: Neural Modal ODEs: Integrating Physics-based Modeling with Neural ODEs for Modeling High Dimensional Monitored Structures. <https://doi.org/10.48550/arXiv.2207.07883>
 - [28] Yildiz, C., Heinonen, M., Lahdesmaki, H.: In: *Advances in Neural Information Processing Systems*
 - [29] Du, J., Futoma, J., Doshi-Velez, F.: Model-based reinforcement learning for semi-markov decision processes with neural odes. NIPS '20. Curran Associates Inc., Red Hook, NY, USA (2020)
 - [30] Paoletti, M.E., Haut, J.M., Plaza, J., Plaza, A.: Neural ordinary differential equations for hyperspectral image classification. *IEEE Transactions on Geoscience and Remote Sensing* **58**(3), 1718–1734 (2020) <https://doi.org/10.1109/TGRS.2019.2948031>
 - [31] Ruiz-Balet, D., Zuazua, E.: Neural ode control for classification, approximation, and transport. *SIAM Review* **65**(3), 735–773 (2023) <https://doi.org/10.1137/21M1411433> <https://doi.org/10.1137/21M1411433>
 - [32] Hairer, E., Norsett, S., Wanner, G.: *Solving Ordinary Differential Equations I: Nonstiff Problems* vol. 8, (1993). <https://doi.org/10.1007/978-3-540-78862-1>
 - [33] Kingma, D.P., Welling, M.: Auto-encoding variational bayes. *CoRR*

abs/1312.6114 (2013)

- [34] Kidger, P., Morrill, J., Foster, J., Lyons, T.: Neural controlled differential equations for irregular time series. In: Proceedings of the 34th International Conference on Neural Information Processing Systems. NIPS '20. Curran Associates Inc., Red Hook, NY, USA (2020)
- [35] Herrera, C., Krach, F., Teichmann, J.: Neural jump ordinary differential equations: Consistent continuous-time prediction and filtering. arXiv: Machine Learning (2020)
- [36] Seedat, N., Imrie, F., Bellot, A., Qian, Z., Schaar, M.: Continuous-time modeling of counterfactual outcomes using neural controlled differential equations. In: International Conference on Machine Learning (2022). <https://api.semanticscholar.org/CorpusID:249712326>
- [37] Ouala, S., Fablet, R., Drumetz, L., Chapron, B., Pascual, A., Collard, F., Gaultier, L.: Physically informed neural networks for the simulation and data-assimilation of geophysical dynamics, pp. 3490–3493 (2020). <https://doi.org/10.1109/IGARSS39084.2020.9323081>
- [38] Ouala, S., Nguyen, D., Drumetz, L., Chapron, B., Pascual, A., Collard, F., Gaultier, L., Fablet, R.: Learning latent dynamics for partially observed chaotic systems. *Chaos: An Interdisciplinary Journal of Nonlinear Science* **30**, 103121 (2020) <https://doi.org/10.1063/5.0019309>
- [39] Arnold, V.I.: Ordinary Differential Equations. Springer, Berlin, Heidelberg (1992). <https://doi.org/10.1007/978-3-642-58088-8>
- [40] Hairer, E., Wanner, G.: Solving Ordinary Differential Equations II. Stiff and Differential-Algebraic Problems vol. 14, (1996). <https://doi.org/10.1007/>

- [41] Hogg, R.V., McKean, J.W., Craig, A.T.: Introduction to Mathematical Statistics. Pearson education international. Pearson Education, Pearson/Prentice Hall, Upper Saddle River, NJ, 64 (2005)
- [42] Ribeiro, M., Ribeiro, I.: Kalman and Extended Kalman Filters: Concept, Derivation and Properties
- [43] Julier, S.J., Uhlmann, J.K.: New extension of the kalman filter to nonlinear systems. In: Defense, Security, and Sensing (1997). <https://api.semanticscholar.org/CorpusID:7937456>
- [44] Särkkä, S.: Bayesian Filtering and Smoothing, (2013). <https://doi.org/10.1017/CBO9781139344203>
- [45] Hansen, L.P.: Large sample properties of generalized method of moments estimators. *Econometrica* **50**(4), 1029–1054 (1982)
- [46] Li, X., Wong, T.-K.L., Chen, R.T.Q., Duvenaud, D.: Scalable gradients for stochastic differential equations. In: Chiappa, S., Calandra, R. (eds.) Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics. Proceedings of Machine Learning Research, vol. 108, pp. 3870–3882. PMLR, Brookline, MA (2020). <https://proceedings.mlr.press/v108/li20i.html>
- [47] Kidger, P., Foster, J., Li, X.C., Lyons, T.: Efficient and accurate gradients for neural sdes. In: Ranzato, M., Beygelzimer, A., Dauphin, Y., Liang, P.S., Vaughan, J.W. (eds.) Advances in Neural Information Processing Systems, vol. 34, pp. 18747–18761. Curran Associates, Inc., Brookline, MA (2021). https://proceedings.neurips.cc/paper_files/paper/2021/file/9ba196c7a6e89eafd0954de80fc1b224-Paper.pdf

- [48] Malekzadeh, M., Clegg, R.G., Cavallaro, A., Haddadi, H.: Mobile sensor data anonymization. In: Proceedings of the International Conference on Internet of Things Design and Implementation. IoTDI '19, pp. 49–58. ACM, New York, NY, USA (2019). <https://doi.org/10.1145/3302505.3310068> . <http://doi.acm.org/10.1145/3302505.3310068>
- [49] Winter, D.A.: Biomechanics and Motor Control of Human Movement, 4th edn. John Wiley & Sons, Hoboken, New Jersey (2009)
- [50] McGeer, T.: Passive dynamic walking. I. J. Robotic Res. **9**, 62–82 (1990) <https://doi.org/10.1177/0278364990000900206>
- [51] Saha, U., Saha, S., Kabir, M.T., Fattah, S.A., Saquib, M.: Decoding human activities: Analyzing wearable accelerometer and gyroscope data for activity recognition. IEEE Sensors Letters **8**(8), 1–4 (2024) <https://doi.org/10.1109/LSENS.2024.3423340>
- [52] Bodhe, R., Sivakumar, S., Sakarkar, G., Juwono, F., Apriono, C.: Outdoor activity classification using smartphone based inertial sensor measurements. Multimedia Tools and Applications **83**, 1–27 (2024) <https://doi.org/10.1007/s11042-024-18599-w>

Appendix A Proofs and examples

Proof of Theorem 1. First, estimate the conditional probability of misclassification

$$\begin{aligned} \mathbb{P}(\delta(\mathbf{x}(t_i)) = j | y = p) &= \mathbb{P}\left[\bigcap_{s=1, s \neq j}^K \{\rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f_j)) < \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f_s))\} \mid y = p\right] \leq \\ &| \text{monotonicity of probability} | \leq \mathbb{P}[\rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f_j)) < \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i | \bar{\mathbf{x}}_0, f_p)) \mid y = p]. \end{aligned}$$

Expand the metrics in the last event using $y = p$:

$$\rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) = \sum_{i=1}^n \|\mathbf{x}(t_i) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)\|^2 = \sum_{i=1}^n \|\boldsymbol{\epsilon}_{t_i}\|^2,$$

$$\begin{aligned} \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)) &= \sum_{i=1}^n \|\mathbf{x}(t_i) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)\|^2 = \sum_{i=1}^n \|(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)) + \boldsymbol{\epsilon}_{t_i}\|^2 = \\ &= \sum_{i=1}^n \|\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)\|^2 + 2 \sum_{i=1}^n \langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \boldsymbol{\epsilon}_{t_i} \rangle + \sum_{i=1}^n \|\boldsymbol{\epsilon}_{t_i}\|^2. \end{aligned}$$

Here $\langle \cdot, \cdot \rangle$ is the standard scalar product. The last event is now

$$\begin{aligned} \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)) &< \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) \Leftrightarrow \\ \sum_{i=1}^n \langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \boldsymbol{\epsilon}_{t_i} \rangle &< -\frac{1}{2} \sum_{i=1}^n \|\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)\|^2 \Leftrightarrow \\ \sum_{i=1}^n \langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \boldsymbol{\epsilon}_{t_i} \rangle &< -\frac{1}{2} \rho_n(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)). \end{aligned}$$

Apply Chebyshev's inequality and use (3):

$$\begin{aligned} \mathbb{P}[\rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)) < \rho_n(\mathbf{x}(t_i), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) \mid y = p] &\leq 4 \frac{\mathbb{D}\left[\sum_{i=1}^n \langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \boldsymbol{\epsilon}_{t_i} \rangle\right]}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \\ &= 4 \frac{\sum_{i=1}^n \mathbb{E}[\langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \boldsymbol{\epsilon}_{t_i} \rangle^2]}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \\ &= 4 \frac{\sum_{i=1}^n (\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j))^T \mathbb{E}[\boldsymbol{\epsilon}_{t_i} \boldsymbol{\epsilon}_{t_i}^T] (\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j))}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \\ &= 4 \frac{\sum_{i=1}^n (\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j))^T R (\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j))}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \end{aligned}$$

$$\begin{aligned}
&= 4 \frac{\sum_{i=1}^n \langle \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), R(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)) \rangle}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} \leq \\
&\leq 4\sigma^2 \frac{\sum_{i=1}^n \|\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j)\|^2}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \\
&= 4\sigma^2 \frac{\rho_n(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))}{\rho_n^2(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} = \frac{4\sigma^2}{\rho_n(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p))} := \xi_n.
\end{aligned}$$

By the theorem condition we have

$$\rho_n(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) \rightarrow \rho_\infty(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_j), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f_p)) = +\infty, \quad n \rightarrow \infty,$$

so $\xi_n \rightarrow 0$, $n \rightarrow \infty$ for all $j, p \in 1 \dots K$.

Second, rewrite the probability of misclassification via conditional probabilities:

$$\begin{aligned}
\mathbb{P}(\delta(\mathbf{x}(t_i)) \neq y) &= \sum_{p=1}^K \mathbb{P}(y = p) \mathbb{P}(\delta(\mathbf{x}(t_i)) \neq p | y = p) \leq \sum_{p=1}^K \mathbb{P}\left(\bigcup_{j=1, j \neq p}^K \{\delta(\mathbf{x}(t_i)) = j\} \mid y = p\right) \leq \\
| \text{union bound} | &\leq \sum_{p=1}^K \sum_{j=1, j \neq p}^K \mathbb{P}(\delta(\mathbf{x}(t_i)) = j | y = p) \leq \sum_{p=1}^K \sum_{j=1, j \neq p}^K \xi_n = K(K-1)\xi_n \rightarrow 0.
\end{aligned}$$

□

Example 1. Here we construct a system of two arbitrarily close vector fields with diverging phase trajectories. Let $M = \{\bar{\mathbf{x}} \in \mathbb{R}^2 \mid \|\bar{\mathbf{x}}\|_2 \leq R\}$ be a disk in the plane. Define the vector field's norm as the sup norm

$$\|f\| = \sup_{\bar{\mathbf{x}} \in M} \|f(\bar{\mathbf{x}})\|_2.$$

Let $\bar{\mathbf{x}}_0 = c \cdot (\sin \gamma, \cos \gamma)^T \neq \mathbf{0}$, $c > 0$, be an initial value. Define the vector fields as

$$\begin{aligned} f(\bar{\mathbf{x}}) &= \begin{pmatrix} 0 & -a \\ a & 0 \end{pmatrix} \bar{\mathbf{x}} = A\bar{\mathbf{x}}, \\ \tilde{f}(\bar{\mathbf{x}}) &= \begin{pmatrix} -b & -a \\ a & -b \end{pmatrix} \bar{\mathbf{x}} = \tilde{A}\bar{\mathbf{x}}, \end{aligned}$$

where $a \neq 0, b > 0$ are some scalars. Then

$$\|f - \tilde{f}\| = \sup_{\bar{\mathbf{x}} \in M} \|(A - \tilde{A})\bar{\mathbf{x}}\|_2 \leq \|A - \tilde{A}\|_2 \cdot R = \left\| \begin{pmatrix} b & 0 \\ 0 & b \end{pmatrix} \right\|_2 \cdot R = bR,$$

where $\|A - \tilde{A}\|_2$ is the matrix spectral norm. So if $b < \frac{r}{R}$ then $\|f - \tilde{f}\| < r$. Next, it is easy to obtain that the IVP solutions for the ODE systems are

$$\begin{aligned} \bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f) &= c \begin{pmatrix} \sin(\gamma - at) \\ \cos(\gamma - at) \end{pmatrix}, \\ \bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, \tilde{f}) &= c \exp(-bt) \begin{pmatrix} \sin(\gamma - at) \\ \cos(\gamma - at) \end{pmatrix}. \end{aligned}$$

It is obvious that $\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, \tilde{f}) \rightarrow \mathbf{0}, t \rightarrow \infty$ and $\|\bar{\mathbf{x}}(t|\bar{\mathbf{x}}_0, f)\| = \|\bar{\mathbf{x}}_0\| = c \forall t$. So

$$\|\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, \tilde{f})\| \rightarrow c \neq 0, \quad i \rightarrow \infty,$$

and then

$$\rho_n(\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f), \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, \tilde{f})) = \sum_{i=1}^n \|\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, f) - \bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, \tilde{f})\|_2^2 \rightarrow \infty, \quad n \rightarrow \infty,$$

because the necessary condition for series convergence does not hold.

As a result, the classifier (7) that uses $\bar{\mathbf{x}}(t_i|\bar{\mathbf{x}}_0, \tilde{f})$ does not have the guarantees of Theorem 1 even though f and $\tilde{f}$ are arbitrarily close.